

1 Revisit the Polynomial Curve Fitting problem

2 Linear Basis Function

- Linear model $y(\mathbf{x}, \mathbf{w}) = w_0 + w_1x_1 + \cdots + w_Dx_D$
- There is an obvious limitation of the linear model (what is it?), we can overcome the limitation by introducing fixed non-linear transformations, using basis functions $\phi_j(\mathbf{x})$.

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x})$$

- Add additional dummy basis $\phi_0(\mathbf{x}) = 1$ to match bias x_0 :

$$y(\mathbf{x}, \mathbf{w}) = \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \Phi(\mathbf{x})$$

- In the curve fitting problem we used $\phi_j(x) = x^j$, polynomial basis functions.
- Applying basis functions can be considered as the feature engineering process, from which we obtain new (and hopefully better) features from original feature space.
- Other examples of basis functions
 - Gaussian basis functions (spline functions) $\phi_j = \exp\left\{-\frac{x-\mu_j}{2s^2}\right\}$
 - Sigmoidal basis function $\phi_j(x) = \sigma\left(\frac{x-\mu_j}{s}\right)$, where $\sigma(a) = \frac{1}{1+\exp(-a)}$
 - Fourier basis (Fourier transformation)

3 From Gaussian noise to Least Squares

- Assumption: observation from a deterministic function with added Gaussian noise:

$$t = y(\mathbf{x}; \mathbf{w}) + \epsilon \quad \text{where} \quad p(\epsilon|\beta) = \mathcal{N}(\epsilon|0, \beta^{-1})$$

- Equivalently: $p(t|\mathbf{x}, \mathbf{w}, \beta) = \mathcal{N}(t - y(\mathbf{x}; \mathbf{w})|0, \beta^{-1}) = \mathcal{N}(t|y(\mathbf{x}; \mathbf{w}), \beta^{-1})$
- Given a dataset $\mathcal{D} = \{\mathbf{X}, \mathbf{t}\}$, where $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ are features, and $\mathbf{t} = \{t_1, \dots, t_N\}$ are targets.
- When samples are drawn from *i.i.d.*, the probability of data is given by:

$$p(\mathcal{D}; \mathbf{w}, \beta) = p(\mathbf{X}, \mathbf{t}; \mathbf{w}, \beta) = p(\mathbf{t}|\mathbf{X}; \mathbf{w}, \beta)p(\mathbf{X})$$

- With the linear prediction function $y(\mathbf{x}_n; \mathbf{w}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n)$, the likelihood function is given by:

$$p(\mathbf{t}|\mathbf{X}; \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n|y(\mathbf{x}_n; \mathbf{w}), \beta^{-1}) = \prod_{n=1}^N \mathcal{N}(t_n|\mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n), \beta^{-1}).$$

- And the log-likelihood is:

$$\begin{aligned} \ln p(\mathbf{t}|\mathbf{X}; \mathbf{w}, \beta) &= \ln \prod_{n=1}^N \mathcal{N}(t_n|\mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n), \beta^{-1}) = \sum_{n=1}^N \ln \mathcal{N}(t_n|\mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n), \beta^{-1}) \\ &= -\frac{N}{2} \ln(2\pi) + \frac{N}{2} \ln \beta - \beta \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2 \end{aligned}$$

- Our goal is to maximize the log-likelihood. Removing constants, we get the following:

$$\max_{\mathbf{w}, \beta} \ln p(\mathbf{t}|\mathbf{X}; \mathbf{w}, \beta) = \min_{\mathbf{w}, \beta} \frac{\beta}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2 - \frac{N}{2} \ln \beta$$

We note the problem of \mathbf{w} is independent from β :

$$\mathbf{w}_{\text{ML}} = \underset{\mathbf{w}}{\text{argmin}} \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2$$

Plugin in \mathbf{w}_{ML} , the optimal solution of the inverse of noise precision β is given by:

$$\frac{1}{\beta_{\text{ML}}} = \frac{1}{N} \sum_{n=1}^N \{t_n - \mathbf{w}_{\text{ML}}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2,$$

which is the residual variance of the target values around the regression function.

- Introducing design matrix $\boldsymbol{\Phi} \in \mathbb{R}^{N \times M}$, N samples and M features (after feature engineering):

$$\boldsymbol{\Phi} = \begin{pmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \dots & \phi_{M-1}(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & \dots & \phi_{M-1}(\mathbf{x}_2) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}_N) & \phi_1(\mathbf{x}_N) & \dots & \phi_{M-1}(\mathbf{x}_N) \end{pmatrix}$$

The least squares can then be represented as:

$$\min_{\mathbf{w}} \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n)\}^2 = \min_{\mathbf{w}} \left\{ \frac{1}{2} \|\boldsymbol{\Phi} \mathbf{w} - \mathbf{t}\|_2^2 = \nabla E_D \right\}.$$

4 Solving Least Squares

4.1 Unregularized Version

- Solve least squares problem using SVD: $\min_{\mathbf{w} \in \mathbb{R}^n} \|\boldsymbol{\Phi} \mathbf{w} - \mathbf{t}\|_2^2$.

- Assume the SVD of Φ is given by

$$\Phi = (U_1 \ U_2) \begin{pmatrix} \Sigma \\ 0 \end{pmatrix} V^T,$$

where $\Sigma \in \mathbb{R}^{N \times M}$ is diagonal, $U_1 \in \mathbb{R}^{N \times M}$ and $U_2 \in \mathbb{R}^{N \times (M-N)}$ have orthonormal columns and $V \in \mathbb{R}^{M \times M}$ is orthogonal.

$$\begin{aligned} \|\Phi \mathbf{w} - \mathbf{t}\|_2^2 &= \left\| (U_1 \ U_2) \begin{pmatrix} \Sigma \\ 0 \end{pmatrix} V^T \mathbf{w} - \mathbf{t} \right\|_2^2 = \left\| \begin{pmatrix} \Sigma \\ 0 \end{pmatrix} V^T \mathbf{w} - (U_1 \ U_2)^T \mathbf{t} \right\|_2^2 \\ &= \left\| \begin{pmatrix} \Sigma y \\ 0 \end{pmatrix} - \begin{pmatrix} b_1 \\ b_2 \end{pmatrix} \right\|_2^2 = \|\Sigma y - b_1\|_2^2 + \|b_2\|_2^2, \end{aligned}$$

where $y = V^T \mathbf{w}$, $b_1 = U_1^T \mathbf{t}$, and $b_2 = U_2^T \mathbf{t}$. The optimal y is given by $y = \Sigma^{-1} b_1$. Thus, the least squares solution is given by

$$\mathbf{w} = Vy = V\Sigma^{-1}b_1 = V\Sigma^{-1}U_1^T \mathbf{t} = \sum_{i=1}^M \frac{u_i^T \mathbf{t}}{\sigma_i} v_i.$$

which is a weighted average of right singular vectors.

- What if $\sigma_i \rightarrow 0$?
- Alternatively solution: set the gradient of least squares to 0

$$\nabla_{\mathbf{w}} \frac{1}{2} \|\Phi \mathbf{w} - \mathbf{t}\|_2^2 = \Phi^T (\Phi \mathbf{w} - \mathbf{t}) = 0 \Rightarrow \mathbf{w}^* = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

- Therefore $\Phi \mathbf{w}^* = \Phi (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$.
- Geometry: This leads to a projection to the space spanned by the columns of Φ (why?)
- When $\Phi^T \Phi$ is close to singular, we would have numerical problems.
- We call $\Phi^\dagger \equiv (\Phi^T \Phi)^{-1} \Phi^T$ pseudo-inverse (generalization of inverse to non-square matrix). See for a square invertible matrix Φ , we have $\Phi^\dagger = (\Phi^T \Phi)^{-1} \Phi^T = \Phi^{-1} (\Phi^T)^{-1} \Phi^T = \Phi^{-1}$
- When $\text{rank}(\Phi) = r \leq \min(M, N)$, let $\Phi = U_r \Sigma_r V_r^T$ be the economical SVD

$$\mathbf{w} = (V_r \Sigma_r U_r^T U_r \Sigma_r V_r^T)^{-1} V_r \Sigma_r U_r^T \mathbf{t} = \sum_{i=1}^r \frac{u_i^T \mathbf{t}}{\sigma_i} v_i$$

- How about when $\sigma_i \rightarrow 0$?
- Since SVD is too expensive, typically we use gradient descent. Conceptually for each iteration:

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \nabla E_D = \mathbf{w}^t - \eta (\Phi^T \Phi \mathbf{w} - \Phi^T \mathbf{t})$$

How to efficient compute the gradient $\Phi^T \Phi \mathbf{w} - \Phi^T \mathbf{t}$?

- Recall one classical problem in dynamic programming – Matrix Chain Multiplication.

$$* \ A \in \mathbb{R}^{10 \times 30}, B \in \mathbb{R}^{30 \times 5}, C \in \mathbb{R}^{5 \times 60}$$

- What if data is very large, and compute the gradient is nearly impossible? We can use stochastic gradient. Each time we see one data point n , we have loss function $E_n = (t_n - \mathbf{w}^{(t)T} \phi_n)^2$, then we perform gradient descent.

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \nabla E_n = \mathbf{w}^t - \eta \phi_n (\mathbf{w}^{(t)T} \phi_n - t_n)$$

where $\phi_n = \phi(\mathbf{x}_n)$

- How about data is located in different servers ($\{\mathcal{S}_1, \mathcal{S}_2, \dots, \mathcal{S}_s\}$)?

$$\nabla E_D = \Phi^T (\Phi \mathbf{w} - \mathbf{t}) = \sum_{i=1}^N \phi_i (\phi_i^T \mathbf{w} - t_i) = \sum_{i \in \mathcal{S}_1} \phi_i (\phi_i^T \mathbf{w} - t_i) + \dots + \sum_{i \in \mathcal{S}_s} \phi_i (\phi_i^T \mathbf{w} - t_i)$$

4.2 Multiple Outputs

- Predict K targets simultaneously $\mathbf{y}(\mathbf{x}, \mathbf{w}) = \mathbf{W}^T \phi(\mathbf{x})$
- Assume the observation is from a multi-response deterministic function with added Isotropic Gaussian distribution $\mathbf{t} = \mathbf{W}^T \phi(\mathbf{x}) + \epsilon$, then $p(\mathbf{t}|\mathbf{x}, \mathbf{W}, \beta) = \mathcal{N}(\mathbf{t}|\mathbf{W}^T \phi(\mathbf{x}), \beta^{-1} \mathbf{I})$.
- Given a set of observations t_1, \dots, t_N , we can combine these into a matrix $\mathbf{T} \in \mathbb{R}^{N \times K}$, and the likelihood function is given by

$$\begin{aligned} \ln p(\mathbf{T}|\mathbf{X}, \mathbf{W}, \beta) &= \sum_{n=1}^N \ln \mathcal{N}(\mathbf{t}_n | \mathbf{W}^T \phi(\mathbf{x}_n), \beta^{-1} \mathbf{I}) \\ &= \frac{NK}{2} \ln \left(\frac{\beta}{2\pi} \right) - \frac{\beta}{2} \sum_{n=1}^N \|\mathbf{t}_n - \mathbf{W}^T \phi(\mathbf{x}_n)\|^2 \end{aligned}$$

- Maximize the log-likelihood gives $\mathbf{W}_{\text{ML}} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{T}$.
- Compare the solution from solving a multi-response least squares, and the solution from solving a set of least squares independently. Are they the same?

4.3 Regularization

- Adding regularization to control overfitting

$$\min_{\mathbf{w}} \frac{1}{2} \|\Phi \mathbf{w} - \mathbf{t}\|_2^2 + \frac{\lambda}{2} \|\mathbf{w}\|_2^2 \quad (1)$$

- Setting gradient to zero, and $\Phi = USV^T = U \begin{bmatrix} \Sigma \\ 0 \end{bmatrix} V$

$$\Phi^T \Phi \mathbf{w} - \Phi^T \mathbf{t} + \lambda \mathbf{w} = 0 \Rightarrow \mathbf{w} = V(\Sigma^2 + \lambda I)^{-1} S U^T \mathbf{t} \Rightarrow \mathbf{w} = \sum_{i=1}^M \frac{\sigma_i u_i^T \mathbf{t}}{\sigma_i^2 + \lambda} v_i$$

Note: V has to be square matrix ($\mathbb{R}^{M \times M}$) in order to be an orthogonal matrix (and thus $I = VV^T$).

- Probabilistic interpretation $p(\mathbf{w}|\mathbf{t}, \mathbf{X}, \beta, \lambda) \propto p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta)p(\mathbf{w}|\lambda)$

$$p(\mathbf{w}, \mathbf{t}|\mathbf{X}, \beta, \lambda) = p(\mathbf{w}|\mathbf{t}, \mathbf{X}, \beta, \lambda)p(\mathbf{t}) = p(\mathbf{t}|\mathbf{X}, \mathbf{w}, \beta)p(\mathbf{w}|\lambda)$$

- A general regularization framework

$$\min_{\mathbf{w}} \frac{1}{2} \|\Phi \mathbf{w} - \mathbf{t}\|_2^2 + \frac{\lambda}{2} \sum_{j=1}^M |\mathbf{w}_j|^q = \min_{\mathbf{w}} \frac{1}{2} \|\Phi \mathbf{w} - \mathbf{t}\|_2^2 + \frac{\lambda}{2} \|\mathbf{w}\|_q^q$$

For ridge regression $q = 2$ and Lasso regression $q = 1$. When $q < 1$ the problem is non-convex (more sparsity)

- Interpretation of regularizations via constrained optimization problem
 - From regularized problems to constrained problems. For convex problems, a regularized problem has an equivalent constrained problem.
 - Sparsity from ℓ_p -norm constrained problems: when $p \leq 0$.

5 Decision Theory for Regression

- Choosing a specific estimate $y(\mathbf{x}) = y(\mathbf{x}; \mathbf{w})$ of the value of t for each input \mathbf{x} , given the squared loss $L(t, y(\mathbf{x})) = \{t - y(\mathbf{x})\}^2$

$$\mathbb{E}[L] = \iint L(t, y(\mathbf{x}))p(\mathbf{x}, t) \, d\mathbf{x} \, dt = \iint \{t - y(\mathbf{x})\}^2 p(\mathbf{x}, t) \, d\mathbf{x} \, dt$$

We will need to find the optimal $y(x)$ using the loss function

$$\min_y \mathbb{E}[L] = \min_y \iint \{t - y(\mathbf{x})\}^2 p(\mathbf{x}, t) \, d\mathbf{x} \, dt$$

To obtain the optimal prediction function $y(x)$, we set

$$\begin{aligned} \frac{\delta \mathbb{E}[L]}{\delta y} &= 2 \int \{y(\mathbf{x}) - t\} p(\mathbf{x}, t) \, dt = 0 \\ \Rightarrow \int y(\mathbf{x}) p(\mathbf{x}, t) \, dt - \int t p(\mathbf{x}, t) \, dt &= 0 \\ \Rightarrow y(\mathbf{x}) p(\mathbf{x}) - \int t p(\mathbf{x}, t) \, dt &= 0 \\ \Rightarrow y(\mathbf{x}) &= \frac{\int t p(\mathbf{x}, t) \, dt}{p(\mathbf{x})} = \mathbb{E}_t[t|\mathbf{x}] \end{aligned}$$

The optimal prediction function given by the squared loss is the conditional expectation of t .